

There is an elephant in the room. . .

Jan Salecker » jsaleck@gwdg.de

2020 M08 04

There is an elephant in the room...



THE AUTHOR OF THE WINDOWS FILE
COPY DIALOG VISITS SOME FRIENDS.



PostgreSQL


Motivation

Motivation


My current R workflow consists of:

- ▶ Local desktop/laptop
 - ▶ Daily work
 - ▶ Development of scripts, analyses
 - ▶ Post-processing of simulation results
 - ▶ Writing papers (RMarkdown)
- ▶ GWDG Cloud server running RStudio Server
 - ▶ Conducting heavy load simulation jobs
 - ▶ Submitting jobs to the GWDG HPC
- ▶ GWDG HPC
 - ▶ Only executes simulations sent from the cloud server

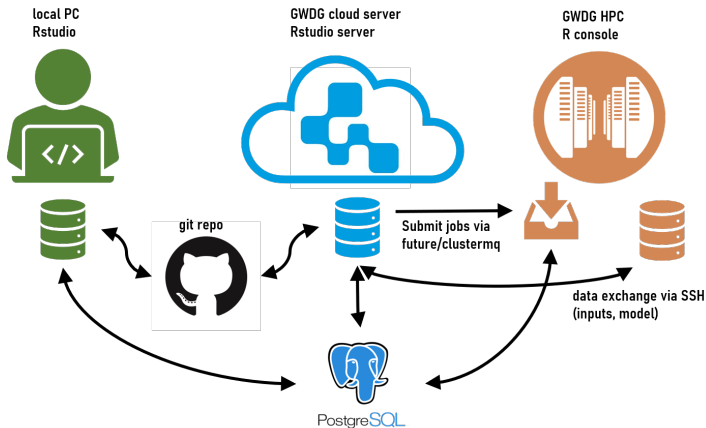
Route to SQL step 1: The Ignorant



Route to SQL step 2: The educated



Route to SQL step 3: The sophisticated



Lessons learned

- ▶ manual copying is bad and error-prone
- ▶ github enables reliable file transfer between local PC and cloud server
- ▶ ssh enables convenient file transfer between cloud server and HPC

But:

- ▶ big data cannot be transferred via github or SSH
- ▶ handling thousands of *.rds files is cumbersome

SQL database is a great addition to manage file transfer between all three instances of R, especially for big data

SQL

- ▶ not as powerful as Apache hadoop but adequate for our purposes
- ▶ database based on tables
 - ▶ relates well to data.frames / tibbles / data.tables
- ▶ R packages allow easy access to SQL databases, even remotely

PostgreSQL

- ▶ open-source version of SQL
- ▶ easy to setup/access
- ▶ unlimited database size & unlimited rows per table
- ▶ maximum table size of 32 TB
- ▶ maximum cell size of 1 GB

How to setup PostgreSQL

How to setup PostgreSQL

I have installed the PostgreSQL server on my GWDG cloud server because it is available 24/7.

I used this detailed guide on how to install the postgresql server:

<https://wiki.ubuntuusers.de/PostgreSQL/>

Afterwards some configuration needs to be done...

Remote access

Step 1: Open SQL communication port

- ▶ Open port 5432 on GWDG cloud server (GWDG services)

Step 2: Edit postgresql.conf file:

- ▶ `sudo nano /etc/postgresql/xx/main/postgresql.conf`
(exchange xx with version number)
- ▶ Change line `listen_addresses='localhost'` to
`listen_addresses='*'`

Remote access

Step 3: Edit hba_file:

In order to find location of hba_file log into postgresql:

- ▶ `sudo -u postgres psql`
- ▶ then type `SHOW hba_file;` and copy the path (e.g. `/etc/postgresql/12/main/pg_hba.conf`)
- ▶ close postgresql: `\q`
- ▶ edit the file: `sudo nano /etc/postgresql/12/main/pg_hba.conf`

Remote access

- ▶ Original settings of pg_hba.conf:

```
# local    replication    all
# host     replication    all          127.0.0.1/32
# host     replication    all          ::1/128
```

- ▶ change replication to all
- ▶ change ip4 host to 0.0.0.0/0
- ▶ change ip6 host to ::0/0
- ▶ Then restart postgresql `sudo /etc/init.d/postgresql restart`

Configure SQL database:

Some important terminal commands:

- ▶ stop/restart/start:
 - ▶ `sudo /etc/init.d/postgresql restart`
 - ▶ `sudo /etc/init.d/postgresql stop`
 - ▶ `sudo /etc/init.d/postgresql start`
- ▶ Create/remove user:
 - ▶ `sudo -u postgres createuser -P -d NUTZERNAME`
 - ▶ `sudo -u postgres dropuser NUTZERNAME`

Configure SQL database:

Some important terminal commands:

- ▶ Create/remove database (each user can have multiple databases):
 - ▶ `sudo -u postgres createdb -O NUTZERNAME DATENBANK`
 - ▶ `sudo -u postgres dropdb DATENBANK`
- ▶ Connect into database locally from terminal:
 - ▶ choose user/database: `psql -d smnws -U NUTZERNAME`
 - ▶ as admin user: `sudo -u postgres psql`

How to use PostgreSQL from R

- ▶ The R packages RPostgres and DBI allow to access SQL databases
- ▶ Assuming we have setup a user account (jan) and a database (projectXY) we can setup a connection:

```
library(RPostgres)
library(DBI)
con <- dbConnect(RPostgres::Postgres(),
                  dbname = 'projectXY',
                  port = 5432,
                  user = 'jan',
                  password = 'xyz')
```

How to use PostgreSQL from R

Print tables in db (each database can have multiple tables)

```
dbListTables(con)
```

Write tables to db:

```
dbWriteTable(con, "mtcars", mtcars)
```

Read tables from db:

```
dbReadTable(con, "mtcars")
```

Modify tables:

```
newlines <- data.frame(mpg=99, cyl=99, disp=99, hp=99, drat=99)
```

```
dbAppendTable(con, "mtcars", newlines)
```

```
dbReadTable(con, "mtcars")
```

Delete tables:

```
dbRemoveTable(con, "mtcars")
```

Disconnect:

Remote access

When remote access is configured correctly (see above) you can enter the remote host in the connection (in my case the IP address of my GWDG cloud server):

```
con <- dbConnect(RPostgres::Postgres(),  
                 dbname = 'projectXY',  
                 host = "141.5.105.225",  
                 port = 5432,  
                 user = 'jan',  
                 password = 'xyz')  
  
dbListTables(con)  
dbDisconnect(con)
```

Write SQL tables in parallel

- ▶ I want to fill up an empty table from the HPC in parallel
- ▶ To do this, the table need to exist in advance
- ▶ Thus, before submitting my jobs I create an empty tibble to initialize the table

```
results_template <- tibble::tibble(name=character(), value=double())  
results_template
```

```
## # A tibble: 0 x 2
```

```
## # ... with 2 variables: name <chr>, value <dbl>
```

```
dbWriteTable(con,  
             "tablename",  
             results_template,  
             overwrite=TRUE)
```

Fill up table in parallel

- We can then use a function that writes directly to that table

```
simfun <- function(x, con, tablename)
{
  # Output (probably another function is called to do some
  out <- tibble::tibble(name=x, value=x)

  # append data
  dbAppendTable(con, "tablename", out)

  # close connection
  dbDisconnect(con)

  # Return nothing
  return()
}
```

Intermission: File tranfser via SSH

File transfer via SSH

```
hpc.upload <- function (from = NA, to = NA, user = NA,
                        host = "transfer.gwdg.de", key = NA)
{
  session <- ssh::ssh_connect(paste0(user, "@", host), key)
  ssh::scp_upload(session, files = from, to = to)
  ssh::ssh_disconnect(session)
}
```

```
hpc.download <- function (from = NA, to = NA, user = NA,
                          host = "transfer.gwdg.de", key = NA)
{
  session <- ssh::ssh_connect(paste0(user, "@", host), key)
  ssh::scp_download(session, files = from, to = to)
  ssh::ssh_disconnect(session)
}
```

File transfer via SSH

```
hpc.delete <- function (folder = NA, user = NA,
                        host = "transfer.gwdg.de", key = NA)
{
  session <- ssh::ssh_connect(paste0(user, "@", host), key)
  ssh::ssh_exec_wait(session,
                      command = paste0("find ",
                                        folder,
                                        "/ -maxdepth 1 -type",
                                        folder),
                      ssh::ssh_disconnect(session)
}
```


A complete workflow example

Introduction

- ▶ I want to use EFForTS-ABM as an example how my workflow looks at the moment
- ▶ Lets imagine I start by implementing a new feature into EFForTS-ABM
- ▶ Afterwards I want to run some analyses on the HPC

Step 01

- ▶ Local PC
 - ▶ Implement features in NetLogo
 - ▶ Prepare analysis scripts
 - ▶ Do some test runs
 - ▶ Push everything to the github repo

Step 02

- ▶ Cloud server
 - ▶ Pull the changes
 - ▶ Send model files and inputs to HPC via SSH
 - ▶ Prepare empty SQL table (dbwriteTable, overwrite=TRUE)
 - ▶ Execute Q function to start HPC jobs
- ▶ Log into HPC terminal to inspect job status
 - ▶ Once all jobs are finished, proceed

Step 03

- ▶ Cloudserver
 - ▶ Load results from SQL table (dbreadtable)
 - ▶ Attach results tibble to nl object
 - ▶ Store complete nl object as *.rds
 - ▶ Push to git repo
- ▶ Local PC
 - ▶ pull git repo
 - ▶ continue analysing simulation results, make plots, write paper,
...

Questions?

Jan Salecker jsaleck@gwdg.de